

SQL to MongoDB Crash Course

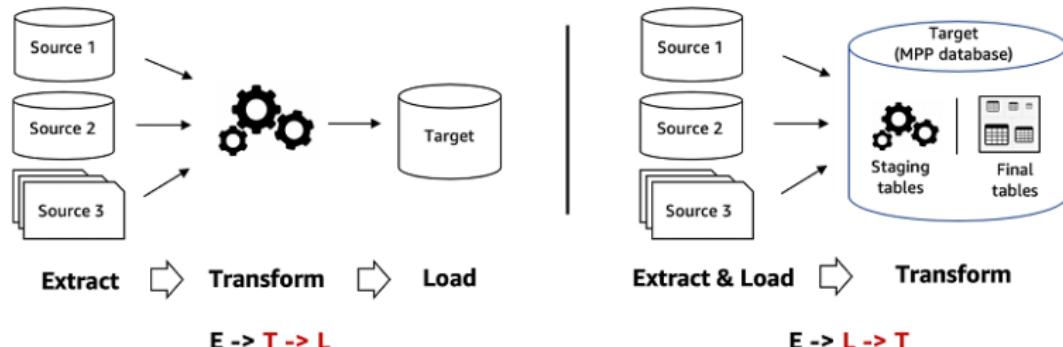
Dr. Sucheta Ghosh

April 10, 2025

Agenda

- Understand the fundamentals of SQL and NoSQL
- Compare SQL and MongoDB
- Explore MongoDB structure and syntax
- Translate SQL queries to MongoDB queries
- Identify real-world simple applications with python SQL/NoSQL

- Data can be numeric (continuous or discrete), categorical (grouped into categories), or ordinal (ordered categories).
- Data can be either structured (organized in a defined format like tables) or unstructured (like free text, images, or videos), and semi-structured.
- ETL vs. ELT

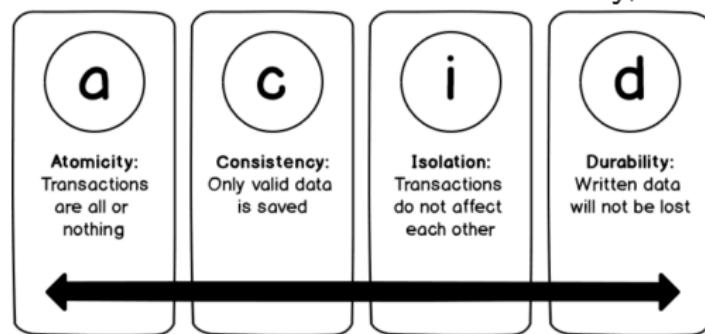


Data & Databases Matter...

- Structured, Semi-Structured, Unstructured data.
- Examples: CSV, JSON, XML, Multimedia.
- Database systems help us organize and query this data.

All started with RDBMS.

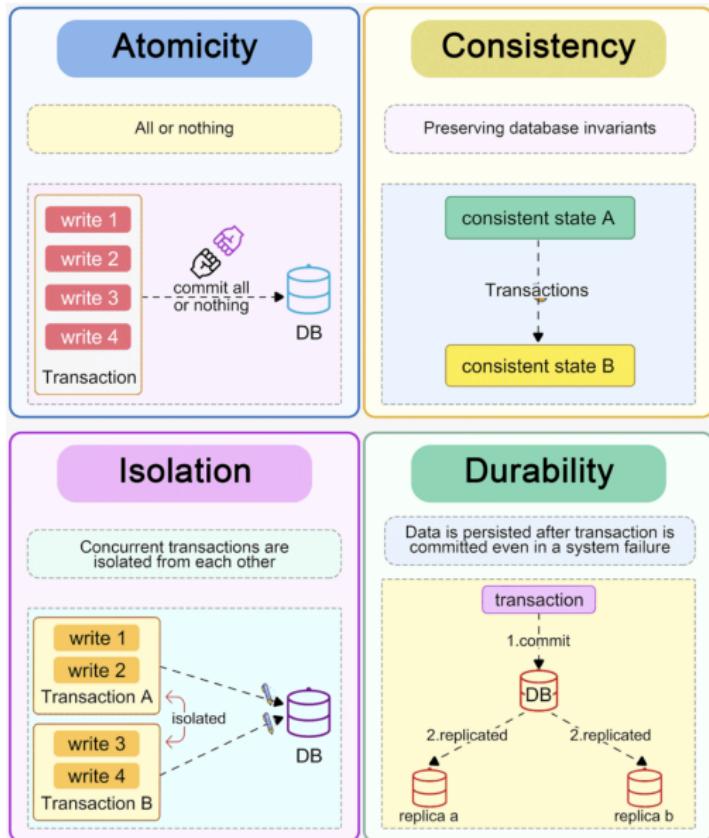
- Data stored in columns and tables
- Relational model with schemas
- Powerful, flexible query language.
- Save Transactions: ACID: Atomicity, Consistency, Isolation, Durability



Note on ACID Properties

- Atomicity: Consider a bank transfer where \$100 is moved from Account A to Account B. If the debit from Account A succeeds but the credit to Account B fails, atomicity ensures that the debit is undone, leaving both accounts unchanged.
- Consistency: In the bank transfer example, if the database has a constraint that an account balance cannot be negative, consistency ensures that no transaction will violate this rule.
- Isolation: If one transaction is transferring \$100 from Account A to Account B and another transaction is reading the balance of Account A, isolation ensures that the second transaction will only see the updated balance after the transfer is fully complete.
- Durability: Once the bank transfer is successfully completed, the changes to the account balances are saved permanently, even if the system crashes immediately after.

ACID Rules



- ACID in short

What is SQL?

- Structured Query Language.
- Works with relational databases.
- Examples: MySQL, PostgreSQL, SQLite.

SQL Structure.

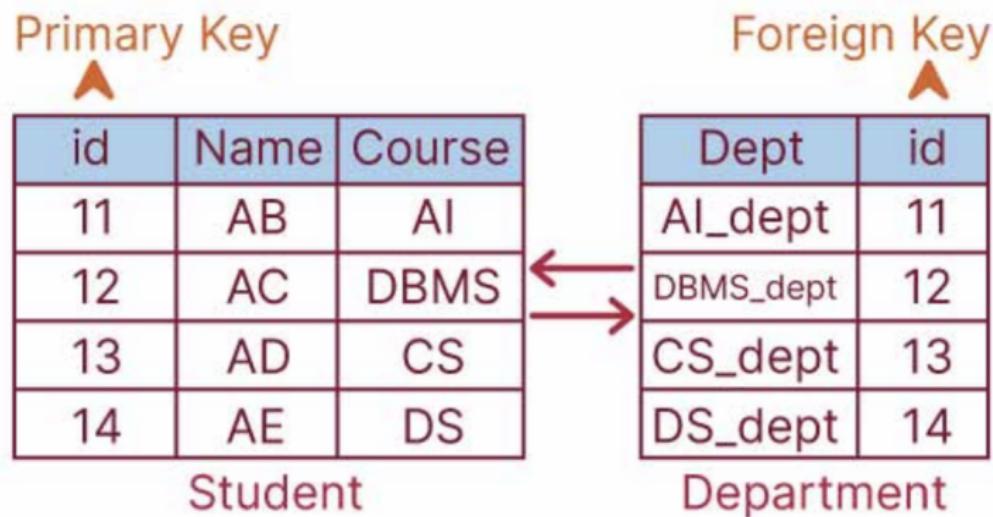
- Tables with rows and columns.
- Fixed schemas.
- Strong consistency.
- Powerful joins

Core SQL Concepts.

- Tables, Primary Keys, Foreign Keys.
- CRUD operations: SELECT, INSERT, UPDATE, DELETE.
- Indexing data on Tables INDEX
- Joins: INNER, LEFT, RIGHT, FULL

SQL: Primary Vs. Secondary Key

- The primary key is a unique identifier within its table, whereas a foreign key is a reference in one table to a primary key in another.
- An example

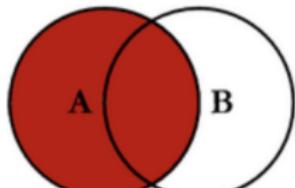


JOINS.

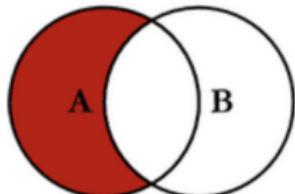
- INNER JOIN: Returns only the rows where there is a match in both tables.
- LEFT JOIN (or LEFT OUTER JOIN): Returns all rows from the left table and matching rows from the right table, filling in NULLs where no match is found.
- RIGHT JOIN (or RIGHT OUTER JOIN): Returns all rows from the right table and matching rows from the left table, filling in NULLs where no match is found.
- FULL JOIN (or FULL OUTER JOIN): Returns all rows from both tables, including all rows from the left table and all rows from the right table, filling in NULLs where no match is found.

JOIN: Set Notation

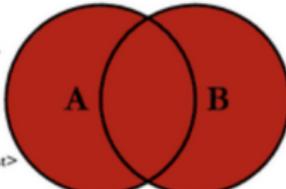
SQL JOINS



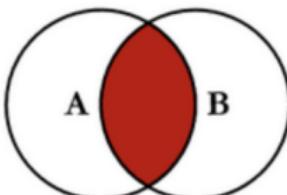
```
SELECT <select_list>
FROM TableA A
LEFT JOIN TableB B
ON A.Key = B.Key
```



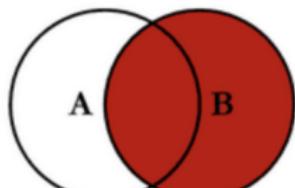
```
SELECT <select_list>
FROM TableA A
LEFT JOIN TableB B
ON A.Key = B.Key
WHERE B.Key IS NULL
```



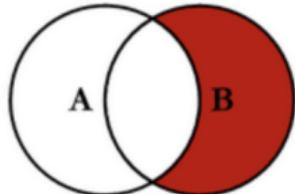
```
SELECT <select_list>
FROM TableA A
FULL OUTER JOIN TableB B
ON A.Key = B.Key
```



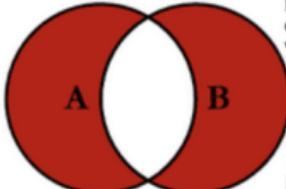
```
SELECT <select_list>
FROM TableA A
INNER JOIN TableB B
ON A.Key = B.Key
```



```
SELECT <select_list>
FROM TableA A
RIGHT JOIN TableB B
ON A.Key = B.Key
```



```
SELECT <select_list>
FROM TableA A
RIGHT JOIN TableB B
ON A.Key = B.Key
WHERE A.Key IS NULL
```



```
SELECT <select_list>
FROM TableA A
FULL OUTER JOIN TableB B
ON A.Key = B.Key
WHERE A.Key IS NULL
OR B.Key IS NULL
```

SQL Syntax Examples.

- `SELECT * FROM orders WHERE price > 100;`
- `SELECT u.name, o.date FROM users u JOIN orders o ON u.id = o.user_id;`

Example: SQL-1.

```
-- Create database
CREATE DATABASE ShopDB;
USE ShopDB;
-- Create tables
CREATE TABLE Customers (
CustomerID INT PRIMARY KEY,
Name VARCHAR(100)
);
CREATE TABLE Orders (
OrderID INT PRIMARY KEY,
CustomerID INT,
Product VARCHAR(100),
FOREIGN KEY (CustomerID) REFERENCES Customers(CustomerID)
);
-- Insert sample data
INSERT INTO Customers VALUES (1, 'Alice'), (2, 'Bob');
INSERT INTO Orders VALUES (101, 1, 'Book'), (102, 2, 'Pen');
```

Example: SQL-2.

```
-- Create index on Name
CREATE INDEX idx_name ON Customers(Name);
-- JOIN Query
SELECT Customers.Name, Orders.Product
FROM Customers
JOIN Orders ON Customers.CustomerID = Orders.CustomerID;
```

SQL Pros and Cons.

- Pros: ACID-compliant, strong schema, rich queries
- Cons: Hard to scale horizontally, rigid schema, JOIN overhead

ACID vs. BASE vs. CAP: ACID

- What it is: A model that emphasizes strong consistency and reliability
- What it stands for: Atomicity, Consistency, Isolation, and Durability
- How it works: Uses a synchronization mechanism to ensure that changes are written to all relevant data sets
- Example: SQL databases use the ACID model.

ACID vs. BASE vs. CAP: BASE

- What it is: A model that prioritizes availability and eventual consistency
- What it stands for: Basically Available, Soft State, and Eventual Consistency
- How it works: Doesn't lock data sets, so it can synchronize without a time guarantee.
- Example: NoSQL databases use the BASE architecture.

ACID vs. BASE vs. CAP: CAP

- What it is: A theorem that highlights the trade-offs between consistency, availability, and partition tolerance.
- What it states: It's not possible to guarantee all three of these properties simultaneously in a distributed system.
- How it applies: A database can satisfy two of the three guarantees.
- Example: The CAP theorem states that it's impossible for a distributed database system to simultaneously guarantee all three properties: consistency, availability, and partition tolerance. Examples of databases that prioritize different combinations of these properties include MongoDB (CP), Cassandra (AP), and DynamoDB (AP).

What is NoSQL?.

- Definition: NoSQL databases are non-relational databases designed for big data and real-time web applications.
- Characteristics:
 - Schema-less data model
 - Horizontal scalability
 - High performance and availability
 - Variety of data models (document, key-value, graph, etc.)

Why NoSQL?.

- Schema flexibility
- Easy to scale horizontally
- Suits modern applications: Web apps, IoT, Logs, Analytics
- Non-relational databases (NoSQL), are needed because they offer advantages in handling large, diverse, and rapidly changing datasets, providing flexibility and scalability that traditional relational databases struggle with.

Meet MongoDB.

- Documents = JSON/BSON
- Collections = Groups of documents
- No fixed schema: dynamic and nested data
- MongoDB = Document-oriented

Overview of MongoDB.

- What is MongoDB?
 - A popular NoSQL document database
 - Stores data in BSON (Binary JSON) format
 - Allows for flexible and dynamic schemas
- Use Cases:
 - Content management systems
 - Real-time analytics
 - Internet of Things (IoT) applications

MongoDB in Action

```
{  
  "name": "Alice",  
  "age": 30,  
  "orders": [  
    { "id": 1, "product": "Pen" },  
    { "id": 2, "product": "Notebook" }  
  ]  
}
```

MongoDB CRUD - Read

```
db.users.find({ age: { $gt: 25 } });
```

MongoDB CRUD - Insert

```
db.users.insertOne({ name: "John", age: 22 });
```

MongoDB CRUD - Insert-Update

```
db.users.updateOne(  
  { name: "John" },  
  { $set: { age: 23 } }  
)
```

MongoDB CRUD - Delete

```
db.users.deleteOne({ name: "John" });
```

Aggregation Framework

- Similar to SQL GROUP BY
- Use '\$match', '\$group', '\$sort'
- Built for analytics queries

Aggregation Example

```
db.sales.aggregate([
  { $match: { price: { $gt: 50 } } },
  { $group: { _id: "$category", total: { $sum: "$price" } } }
]);
```

Indexing in MongoDB

- Improves performance
- Similar to SQL indexes

```
db.users.createIndex( name: 1 );
```

Data Modeling in MongoDB

- Embedded documents vs referenced collections
- Use embedded for 1:1 or 1:many (read-heavy)
- Use reference for normalized (write-heavy)

Embedded Example

```
{ "name": "Alice",  
  "address": {  
    "city": "Berlin",  
    "zip": "10115"  
  }  
}
```

Referenced Example

```
db.users.insertOne( name: "Bob", address_id: ObjectId("...")  
);  
db.addresses.insertOne( _id: ObjectId("..."), city: "Paris" );
```

\$lookup Example (Join)

```
db.orders.aggregate([ $lookup: { from: "products", localField: "product_id", foreignField: "_id", as: "product_details" } ]);
```

SQL vs MongoDB Terminology

SQL	MongoDB
Table	Collection
Row	Document
Column	Field
JOIN	\$lookup or Embedded
Index	Index

SQL to MongoDB Translation Examples

SQL	MongoDB
SELECT * FROM users	db.users.find({})
WHERE age > 30	{ age: \$gt: 30 }
INSERT INTO ...	db.collection.insertOne({...})

CRUD Commands Comparison: SQL vs NoSQL

1. CREATE -- Insert Data

-- SQL

```
INSERT INTO Customers (CustomerID, Name) VALUES (3, 'Charlie');
```

-- NoSQL

```
db.customers.insert_one({"_id": 3, "name": "Charlie"})
```

CRUD Commands Comparison: SQL vs NoSQL

2. READ -- Retrieve Data

-- SQL

-----All records

```
SELECT * FROM Customers;
```

----- Filtered records

```
SELECT * FROM Customers WHERE Name = 'Charlie';
```

-- NoSQL

All documents

```
list(db.customers.find())
```

Filtered query

```
db.customers.find_one("name": "Charlie")
```

CRUD Commands Comparison: SQL vs NoSQL

3. UPDATE -- Modify Data

-- SQL

```
UPDATE Customers SET Name = 'Charles' WHERE CustomerID = 3;
```

-- NoSQL

```
db.customers.update_one(  
  "_id": 3,  
  "$set": "name": "Charles"  
)
```

CRUD Commands Comparison: SQL vs NoSQL

4. DELETE -- Remove Data

-- SQL

```
DELETE FROM Customers WHERE CustomerID = 3;
```

-- NoSQL

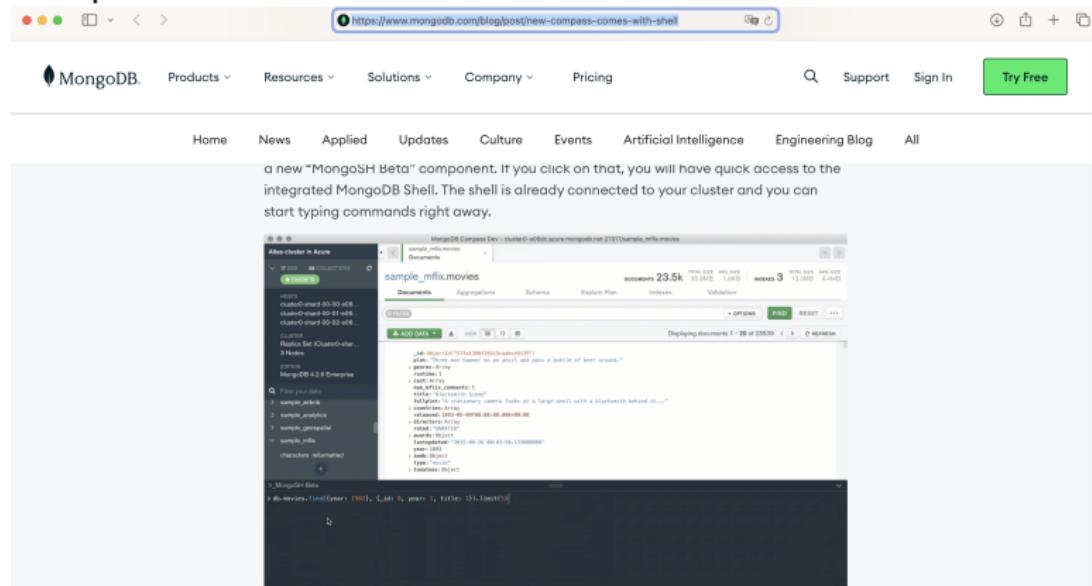
```
db.customers.delete_one({_id": 3})
```

Tooling for MongoDB

- MongoDB Compass (GUI)
- MongoDB Atlas (Cloud DB)
- Mongo Shell / mongosh

MongoDB Compass (Screenshot)

- Browse collections
- Visualize documents
- Build queries and aggregations
- Compass Screenshot



Syntax Highlighting & Autocomplete

The integrated MongoDB Shell includes syntax highlighting to enhance readability and to

SQL to MongoDB Crash Course

MongoDB Atlas (Screenshot)

- Free tier cloud cluster
- Secure and scalable
- Good for prototyping and learning
- Atlas Screenshot

The screenshot shows the MongoDB Atlas web interface. At the top, there's a navigation bar with links for Products, Resources, Solutions, Company, Pricing, Eng (Engineering), Support, Sign In, and a prominent green 'Get Started' button. Below the navigation is a secondary header with tabs for ACME, MYAPP, and Atlas, along with Access Manager, Support, and Billing options. On the left, a sidebar under 'DATA STORAGE' includes 'Clusters' (which is selected and highlighted in green), 'Triggers', 'Data Lake', and sections for 'SECURITY', 'Database Access', 'Network Access', and 'Advanced'. The main content area is titled 'Clusters' and features a large button labeled 'Create a cluster'. Below it, there's a search bar and a note: 'Choose your cloud provider, region, and specs.' A 'Build a Cluster' button is located at the bottom of this section. A note at the bottom right says: 'Once your cluster is up and running, live migrate an existing MongoDB database into Atlas with our Live Migration Service.' At the very bottom of the page, there are links for Feature Requests, System Status (All Good), and footer links for ©2021 MongoDB, Inc., Status, Terms, Privacy, Atlas Blog, Contact Sales.

Limitations of MongoDB

- Lacks strong relational enforcement
- No cross-collection ACID transactions (except v4.0+)
- Indexing and aggregation need tuning for large datasets

SQL + NoSQL Together

- Microservices may use MongoDB
- Reporting layer may use SQL (PostgreSQL/MySQL)
- Polyglot persistence: choose the best DB for the job

Use Cases Comparison

Use Case	SQL	MongoDB
Banking	✓ ACID-safe	X
Product Catalog	X Complex Joins	✓ Flexible
Log Analytics	X	✓ Fast ingest

Recap

- SQL: reliable for structured data
- MongoDB: flexible, fast for dynamic data
- Learn both and apply based on project needs

Quiz: What Would You Use?

- IoT stream data storage?
- Inventory management system?
- Mobile app user profiles?
- (Discuss: SQL or MongoDB and why?)

Thank You!

Questions? sucheta.ghosh@lehrb.hs-offenburg.de